



Creating Core Resources and Building the Diagnostic Ladder



Tim Warner

Principal Author, Pluralsight

TechTrainerTim.com



Four framing questions



How does the Deployment > ReplicaSet > Pod hierarchy enable self-healing workloads?



How do ClusterIP, NodePort, and LoadBalancer Services expose workloads differently?



How do Namespaces, labels, and annotations organize resources at scale?



How does the get > describe > logs > events diagnostic ladder solve every troubleshooting scenario?





“The product catalog API needs to be live by Thursday. Get it running, expose it to the QA team, and make sure you know how to fix it when it goes down. Because it will go down. And troubleshooting is 30% of your CKA score -- so consider this practice for both the job and the exam.”



How does the controller hierarchy enable self-healing workloads?



Bare Pods vs. managed Deployments



Bare Pods

- Created directly with `kubectl run` or YAML
- No controller monitors or replaces failed Pods
- Deleted Pods are gone permanently
- No scaling or self-healing behavior
- Used for debugging, not production workloads



Managed Deployments

- Deployment manages ReplicaSets → Pods
- ReplicaSet continuously enforces desired replica count
- Failed or deleted Pods are automatically recreated
- Supports scaling, rolling updates, and rollbacks
- Standard production pattern for running workloads



The ReplicaSet reconciliation loop

ReplicaSet continuously compares desired count to actual count

If a Pod crashes or is deleted, a replacement is created automatically

If a node fails, Pods are rescheduled to healthy nodes

`kubectl scale deployment` adjusts the desired replica count

You almost never create bare Pods in production or on the exam



How do Services expose workloads to the network?



Services and the selector-label contract



A Service provides a stable virtual IP for a set of Pods



The Service selects Pods by matching their labels



Pod IPs appear in EndpointSlices when labels match



If a Pod loses its label or fails readiness, its IP is removed




EndpointSlices are the modern replacement for the legacy Endpoints object



If labels match and readiness passes, Pod IPs are included in EndpointSlices



K8S request path: DNS → Service → Pod



get Current state

describe Pod events

logs Container output

events Cluster timeline

Follow this ladder to diagnose every troubleshooting scenario.



Three Service types



ClusterIP (default)

Internal-only virtual IP for service-to-service traffic

Stable endpoint for Pods selected by labels

Used by other Pods inside the cluster

DNS: `service-name.namespace.svc.cluster.local`



NodePort

Extends ClusterIP with a port on every node (30000–32767)

Accessible from outside the cluster via node IP + port

Routes traffic to the Service → then to Pods

Works in local clusters (kind, minikube)



LoadBalancer

Extends NodePort with a cloud provider load balancer


Provides external IP for production access

Automatically routes traffic to healthy Pods

On kind: stays Pending (no cloud provider)



Verifying service connectivity with DNS



Every Service gets a DNS record: `service-name.namespace.svc.cluster.local`

Same-namespace Pods can use the short name: just 'my-service'

Short names work via DNS search domains

Cross-namespace access requires the full DNS name

Run from inside a Pod (not your host):

```
kubectl run debug --image=busybox --rm  
-it --nslookup my-service
```

```
kubectl run debug --image=busybox --rm  
-it -- wget -qO- my-service
```



How do Namespaces, labels, and annotations organize resources?



Namespaces, labels, and annotations

Namespaces scope resource names and enable RBAC policies

Labels are for selection
-- Services, ReplicaSets, and NetworkPolicies use them

Annotations carry metadata but do not affect scheduling or routing

`kubectl get pods -A`
queries across all namespaces

Compound selectors:
`kubectl get pods -l env=staging,team=back end`



**How does the diagnostic ladder solve
every troubleshooting scenario?**



The diagnostic ladder

get →
state

describe
→ events

logs →
container

events →
cluster
timeline

Current status (Running, Pending, CrashLoopBackOff)

Scheduling, image pulls, probe failures

stdout / stderr from the container

Chronological, cross-resource events



Container output streams and multi-container Pods

`kubectl logs` reads
stdout and stderr
captured by containerd

For multi-container
Pods: `kubectl logs`
`pod-name -c cont-`
`name`

`kubectl logs --`
`previous` shows
output from the last
terminated container

`kubectl logs --follow`
streams output in real
time

Container output
management is a
dedicated CKA
competency



kubectl logs and multi-container pods

Get logs from a single-container Pod

```
kubectl logs my-pod
```

Multi-container Pod: MUST specify container

```
kubectl logs my-pod -c app-container
```

Previous container instance (CrashLoopBackOff)

```
kubectl logs my-pod --previous
```

Stream logs in real time (live debugging)


```
kubectl logs -f my-pod
```

List containers in a Pod (before choosing -c)

```
kubectl describe pod my-pod
```



Diagnostic ladder applied to the request path



get Current state

describe Pod events

logs Container output

events Cluster timeline

Follow this ladder to diagnose every troubleshooting scenario.



Demo

Deploy, expose, break, and diagnose





“The catalog API went down Friday night. I ran get, describe, logs, events -- found the bad image tag in under two minutes. Three weeks ago that would have been a two-hour panic. The diagnostic ladder works.”





From Globomantics to you



Deploy with Deployments, not bare Pods

Self-healing comes from the ReplicaSet reconciliation loop -- bare Pods do not recover



Verify connectivity end-to-end

DNS resolution plus EndpointSlice inspection proves the full path from client to Pod



Organize with labels, not just namespaces

Labels drive selection for Services, ReplicaSets, and NetworkPolicies -- they are the wiring



Make the diagnostic ladder automatic

get, describe, logs, events -- this pattern is your weapon for 30% of the CKA score





Tim Warner | TechTrainerTim.com

